

Project2-HarvardX-Capstone-Course.R

alevcieslinski

2022-03-09

```
#An Introduction to Topic Modeling and Sentiment Analysis:
#Amazon Food Reviews Data
#This analysis is based on product reviews from Amazon. The original dataset is
#quite large with the following statistics:
#Number of reviews 568,454
#Number of users 256,059
#Number of products 74,258
#I ran my analysis on the whole dataset, but R was very slow, so I decided to
#sample by productid and use the smaller dataset to complete this project. I
# also brought in a subset of the original dataset to make R coding more
#efficient.
#The purpose of this analysis is to do text analytics and topic modeling on a
#sample of Amazon product reviews, mainly food reviews. I also conducted
#sentiment analysis.

#Install and load packages into R
library(dplyr)

##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

library(ggplot2)
library(tidytext)
library(tidyverse) # metapackage of all tidyverse packages

## -- Attaching packages ----- tidyverse 1.3.1 --

## v tibble 3.1.6      v purrr 0.3.4
## v tidyr 1.2.0       v stringr 1.4.0
## v readr 2.1.2       v forcats 0.5.1

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()

library(lubridate)

##
```

```

## Attaching package: 'lubridate'
## The following objects are masked from 'package:base':
##
##     date, intersect, setdiff, union
library(lda)
library(topicmodels)
library(tm)

## Loading required package: NLP
##
## Attaching package: 'NLP'
## The following object is masked from 'package:ggplot2':
##
##     annotate
library(wordcloud)

## Loading required package: RColorBrewer
library(widyr)
library(Matrix)

##
## Attaching package: 'Matrix'
## The following objects are masked from 'package:tidyr':
##
##     expand, pack, unpack
library(kableExtra)

##
## Attaching package: 'kableExtra'
## The following object is masked from 'package:dplyr':
##
##     group_rows
library(textdata)
library(syuzhet)
library(readxl)
library(latexpdf)
library(tinytex)
theme_set(theme_light())

#load dataset into R and eliminate null reviews, assign lower cases to variables
reviews <- read_excel("/Users/alevcieslinski/Downloads/reviews3.xls")
reviews <- na.omit(reviews)
reviews <- reviews %>%
  rename_with(str_to_lower)

#Check the structure of data
str(reviews)

## tibble [5,018 x 3] (S3: tbl_df/tbl/data.frame)
## $ productid: chr [1:5018] "6641040" "6641040" "6641040" "6641040" ...

```

```
## $ text      : chr [1:5018] "A charming, rhyming book that describes the circumstances under which you
## $ score     : num [1:5018] 4 4 5 5 4 5 5 1 4 5 ...
```

```
#data cleanup
```

```
reviews$text <- gsub("<[^\>]+>", "", reviews$text)
```

```
#summary data. I brought in productid, text(review), and score that customers
#assigned products.
```

```
reviews %>% summarize(unique_products = length(unique(productid)),
                      unique_text = length(unique(text)),
                      unique_score = length(unique(score)))
```

```
## # A tibble: 1 x 3
##   unique_products unique_text unique_score
##           <int>       <int>       <int>
## 1             700         4107           5
```

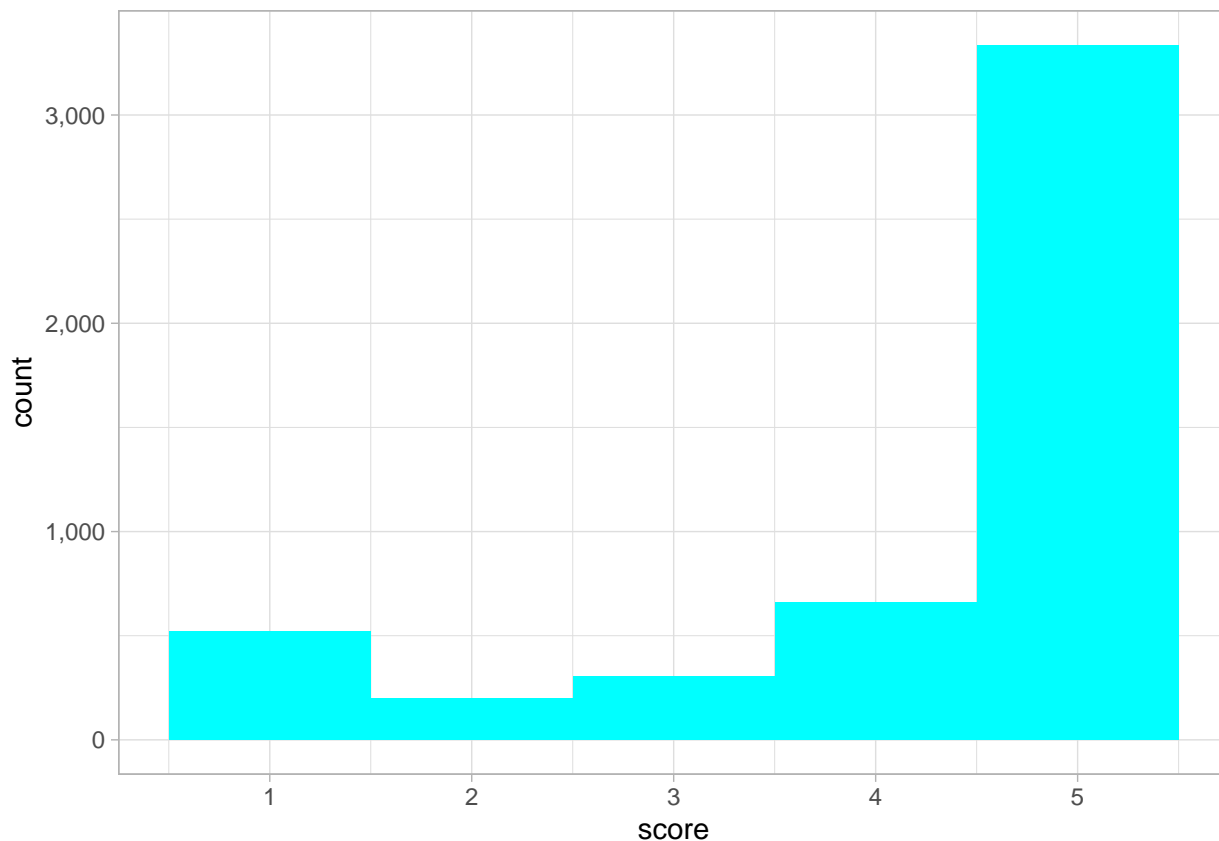
```
#count by score (customer review score)
```

```
reviews %>%
  count(score, sort=TRUE)
```

```
## # A tibble: 5 x 2
##   score     n
##   <dbl> <int>
## 1     5 3336
## 2     4  661
## 3     1  520
## 4     3  304
## 5     2  197
```

```
#visual representation of score distribution
```

```
reviews %>%
  ggplot(aes(score)) +
  geom_histogram(fill="cyan", binwidth=1) +
  scale_y_continuous(labels = scales::comma)
```



```
#tokenization, stopwords
#count words by score and productid
#In order to do the text analysis and topic modeling, we have to create words
#from each line of text.
```

```
food_reviews <- reviews %>%
  unnest_tokens(input=text, output=word) %>%
  anti_join(stop_words, by='word') %>%
  filter(str_detect(word, "[a-z]")) %>%
  count(productid, word, score, sort=TRUE) %>%
  ungroup()
```

```
total_words <- food_reviews %>%
  group_by(productid) %>%
  summarize(total = sum(n))
```

```
food_reviews <- left_join(food_reviews, total_words)
```

```
## Joining, by = "productid"
```

```
#tf_idf
#The statistic tf-idf is intended to measure how important a word is to
#a document in a collection (or corpus) of documents. In this case, we're
#looking at words by productid.
```

```
food_reviews <- food_reviews %>%
  bind_tf_idf(word, productid, n) %>%
  arrange(desc(tf_idf))
glimpse(food_reviews)
```

```
## Rows: 97,158
## Columns: 8
## $ productid <chr> "B0000DD8RB", "B0000DD8RB", "B0000DG56I", "B0000D956L", "B00~
## $ word <chr> "popper", "spits", "spicy", "candies", "stomach.so", "valle~
## $ score <dbl> 3, 3, 4, 4, 2, 5, 5, 5, 5, 5, 1, 1, 4, 1, 5, 5, 5, 5, 5, 5, ~
## $ n <int> 1, 1, 1, 2, 1, 1, 2, 1, 1, 1, 1, 1, 1, 9, 1, 1, 1, 1, 1, 1, ~
## $ total <int> 4, 4, 4, 5, 5, 5, 11, 4, 5, 3, 6, 6, 5, 55, 5, 7, 7, 7, 7, 7~
## $ tf <dbl> 0.2500000, 0.2500000, 0.2500000, 0.4000000, 0.2000000, 0.200~
## $ idf <dbl> 5.857933, 5.857933, 5.857933, 3.660709, 6.551080, 6.551080, ~
## $ tf_idf <dbl> 1.4644833, 1.4644833, 1.4644833, 1.4642834, 1.3102161, 1.310~
```

*#The LDA() function operates on a Document Term Matrix (DTM), so we need to
#create a DTM from the reviews. The dataframe food_reviews containing all of
#the word counts by productid will then be cast into a document term matrix with
#the following code.*

```
dtm <- food_reviews %>%
  cast_dtm(document=productid, term=word, value=n)
dtm
```

```
## <<DocumentTermMatrix (documents: 700, terms: 16790)>>
## Non-/sparse entries: 82193/11670807
## Sparsity : 99%
## Maximal term length: 29
## Weighting : term frequency (tf)
```

*#When we do topic modeling, we need to set some parameters, i.e., the number of
#topics to find in the corpus. For simplicity, we chose 5 topics. LDA is a
#clustering algorithm.*

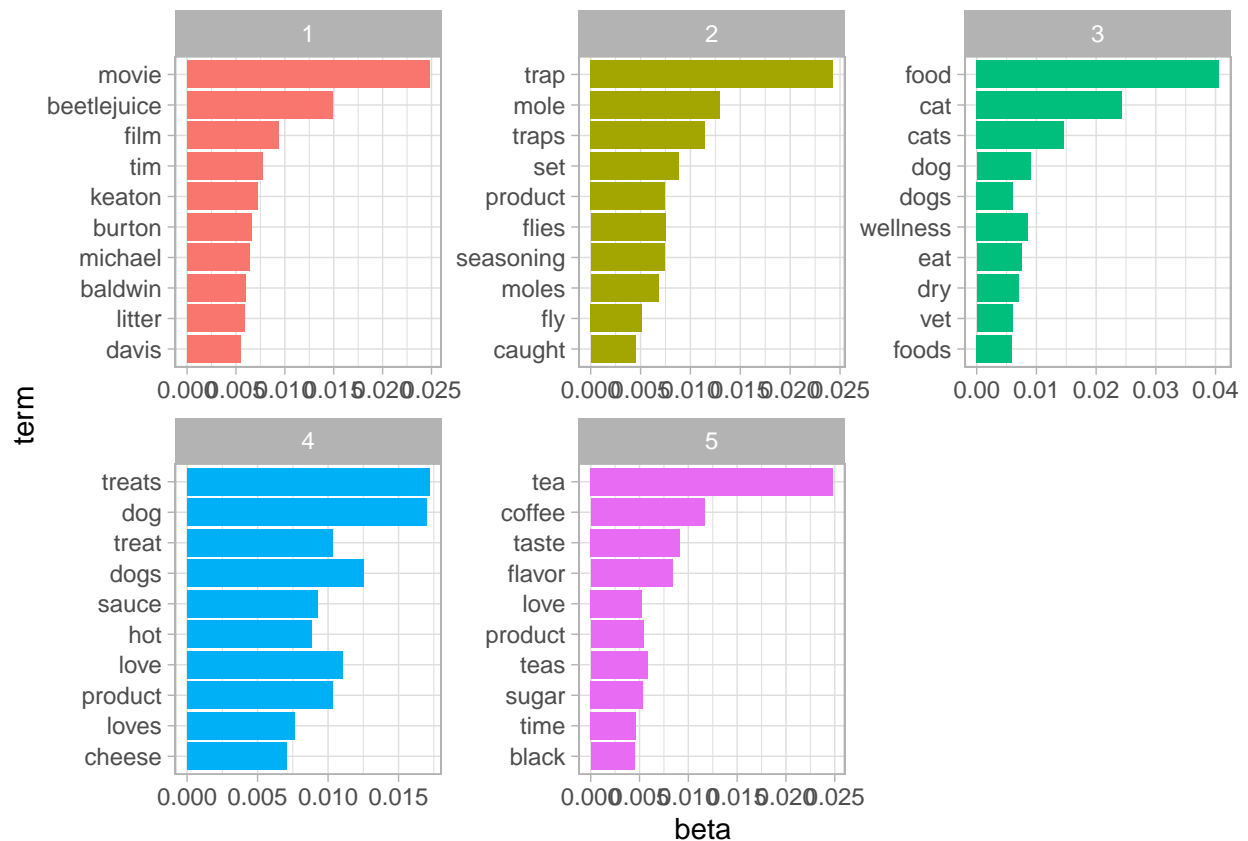
```
lda <- LDA(dtm, k = 5, control = list(seed = 1968))
lda
```

```
## A LDA_VEM topic model with 5 topics.
```

*#We can look at which words occur most commonly within each topic. Beta is
#equal to $P(\text{word}|\text{topic})$, the conditional probability that the word is observed,
#given the topic. The code below look at the top 10 words for each of the five
#topics modeled. Even though the reviews are for food items, there are some
#records with other product reviews. Interestingly, topic modeling was able to
#detect reviews about movies and identify those texts as a distinct set of topic
#and words as shown by the first topic in the plot below.*

```
review_topics <- tidy(lda, matrix = "beta")
review_top_terms <- review_topics %>%
  group_by(topic) %>%
  top_n(10, beta) %>%
  ungroup() %>%
  arrange(topic, -beta)

review_top_terms %>%
  mutate(term = reorder(term, beta)) %>%
  ggplot(aes(term, beta, fill = factor(topic))) +
  geom_col(show.legend = FALSE) +
  facet_wrap(~ topic, scales = "free") +
  coord_flip()
```



*#The product id for the movie reviews is B00004CI84 and we can see the original
#texts for this productid below for further validation:*

```
reviews %>% filter(reviews$productid == "B00004CI84")
```

```
## # A tibble: 189 x 3
##   productid text                                     score
##   <chr>      <chr>                                     <dbl>
## 1 B00004CI84 "... this little gem was the beginning of a brief \"thing\"~    4
## 2 B00004CI84 "\"It keeps getting funnier every time I see it,\" says Be~    5
## 3 B00004CI84 "\"Pee-Wee's Big Adventure\" was director Tim Burton's firs~    5
## 4 B00004CI84 "[[ASIN:B001AGXEAG Beetlejuice (20th Anniversary Deluxe Edi~    4
## 5 B00004CI84 "*** 1/2 stars rating for &quot;Beetlejuice&quot;. This mov~    3
## 6 B00004CI84 "Beetlejuice (20th Anniversary Deluxe Edition)I am very imp~    4
## 7 B00004CI84 "A cool movie with some very funny and amusing moments, a y~    4
## 8 B00004CI84 "a couple dies.they live in a house as spirits and then som~    4
## 9 B00004CI84 "A dead couple attempts to scare a modern family out of the~    5
## 10 B00004CI84 "A haunting movie in which everything is played for laughs.~    1
## # ... with 179 more rows
```

#summarize betas below

```
summary(review_topics)
```

```
##   topic      term      beta
## Min.   :1 Length:83950 Min.   :0.000e+00
## 1st Qu.:2 Class :character 1st Qu.:0.000e+00
## Median :3 Mode  :character Median :0.000e+00
## Mean   :3              Mean  :5.956e-05
```

```
## 3rd Qu.:4          3rd Qu.:3.348e-05
## Max. :5           Max. :4.057e-02
```

*#While beta refers to the conditional probability of a word given a topic,
#gamma is the per-document likelihood of each topic. In the code below,
#for example, we see how the algorithm allocates topics across the first
#two productids.*

```
review_documents <- tidy(lda, matrix = "gamma")
review_documents <- arrange(review_documents, document)
review_documents
```

```
## # A tibble: 3,500 x 3
##   document  topic  gamma
##   <chr>    <int>  <dbl>
## 1 141278509X      1 0.00134
## 2 141278509X      2 0.00134
## 3 141278509X      3 0.00134
## 4 141278509X      4 0.00134
## 5 141278509X      5 0.995
## 6 2734888454      1 0.00159
## 7 2734888454      2 0.00159
## 8 2734888454      3 0.00159
## 9 2734888454      4 0.994
## 10 2734888454      5 0.00159
## # ... with 3,490 more rows
```

Compute perplexity score
#Perplexity score is a measure of how well the model predicts a sample. The
#lower the score, the better the model is deemed to be. Of course, 1 perplexity
#score is not that meaningful as we need a distribution of benchmarks to
#be able to determine the lowest score.

```
perplexity(object=lda, newdata=dtm)
```

```
## [1] 1896.082
```

#Finding the optimal number of k scores
#create a dataframe to store the perplexity scores for different values of k
#As sated above, we create a set of perplexity scores based on different values
#of k to determine the best model.

```
p = data.frame(k = c(2,4,8,16,32,64,128), perplexity = NA)
```

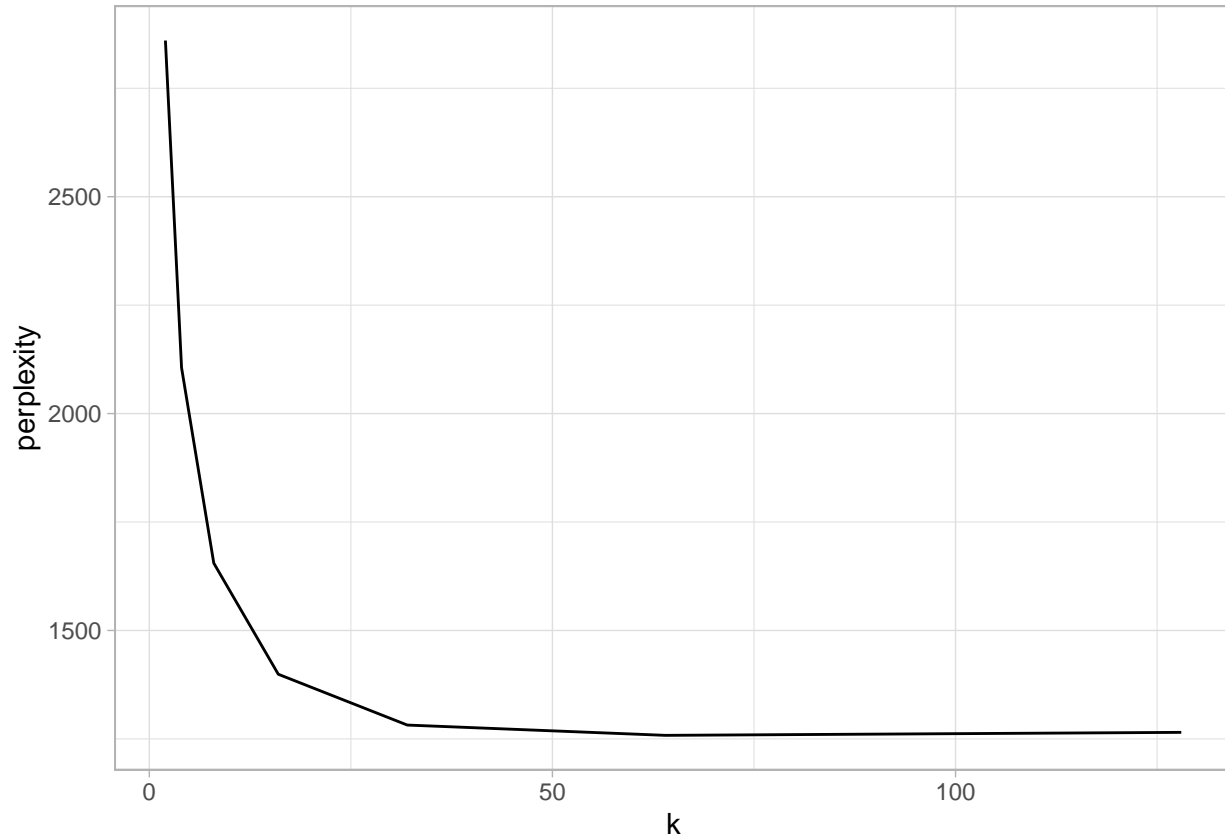
```
# loop over the values of k in data.frame p
for (i in 1:nrow(p)) {
  print(p$k[i])
  #calculate perplexity for the given value of k
  m = LDA(dtm, method = "Gibbs", k = p$k[i], control = list(alpha = 0.01))
  # store result in our data.frame
  p$perplexity[i] = perplexity(m, dtm)
}
```

```
## [1] 2
## [1] 4
## [1] 8
## [1] 16
```

```
## [1] 32
## [1] 64
## [1] 128
```

```
#plot perplexity & values of k
```

```
ggplot(p, aes(x=k, y=perplexity)) + geom_line()
```



```
#we see that the perplexity score sharply declines for values of k greater  
#than 25.
```

```
#wordcloud
```

```
#Next, we can create a wordcloud to see a visual depiction of relative  
#importance of words by generating the counts of words in the corpus
```

```
word_frequencies <- food_reviews %>%  
  count(word)
```

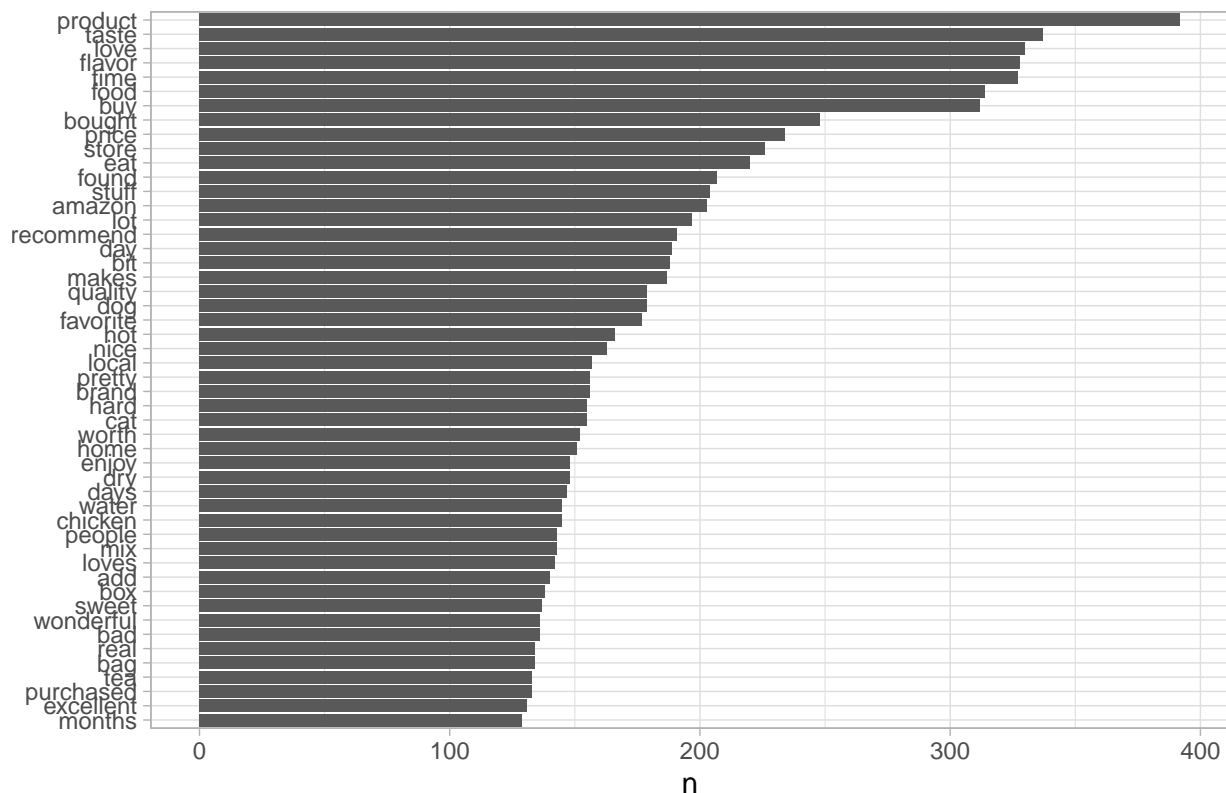
```
#Create the wordcloud
```

```
wordcloud(words=word_frequencies$word,  
          freq=word_frequencies$n,  
          min.freq=5,  
          max.words=50,  
          colors=c("DarkOrange", "Blue"),  
          scale=c(3,0.3))
```




```
#top 50 words that were used on product reviews
food_reviews %>%
  count(word, sort = TRUE) %>%
  head(50) %>%
  mutate(word = reorder(word, n)) %>%
  ggplot(aes(word, n)) +
  geom_col() +
  xlab(NULL) +
  labs(title="Top 50 words used in Reviews") +
  coord_flip()
```

Top 50 words used in Reviews



*#While the wordcloud and the top 50 words graph is helpful in understanding
#what seems to be most used "word/s", it doesn't help us understand
#sentiments and emotions, which is what we will explore next.*

#Sentiment analysis

*#Most words by themselves do not describe sentiment. I used the sentiment
#function below to assign sentiments to the text from the "reviews" dataset.
#Once we run the code and look at s created by the get_nrc_sentiment , we see
#that each row represents a review and each column represents the different
#sentiments along with positive and negative score for that review. We also
#created a final review score that is represented below:*

```
sentiment_data_all <- iconv(reviews$text)
s <- get_nrc_sentiment(sentiment_data_all)
```

```
## Warning: `spread()` was deprecated in tidyr 1.2.0.
## Please use `spread()` instead.
## This warning is displayed once every 8 hours.
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was generated.
```

```
s$score <- s$positive - s$negative
head(sentiment_data_all)
```

```
## [1] "A charming, rhyming book that describes the circumstances under which you eat (or don't) chicken
## [2] "A very entertaining rhyming story--cleaver and catchy.The illustrations are imaginative and fit
## [3] "All of my children love this book. My first grader got it for Christmas and loves to read the
## [4] "Classic children's book, can't go wrong. I read it when i was a kid and ordered it 20yrs later.
## [5] "Get the movie or sound track and sing along with Carol King. This is great stuff, my whole exte
```

```
## [6] "Great book, perfect condition arrived in a short amount of time, long before the expected delivery"
s[1:10,]
```

```
##      anger anticipation disgust fear joy sadness surprise trust negative positive
## 1         1           1       0   2   2         2         0     2         1         6
## 2         0           2       0   0   1         0         0     0         1         2
## 3         0           2       1   1   4         1         1     5         1         9
## 4         0           0       0   0   0         0         0     0         1         1
## 5         0           2       0   0   3         2         0     2         0         4
## 6         0           5       0   0   1         0         0     1         0         2
## 7         0           1       0   0   1         0         0     3         0         1
## 8         1           5       1   0   2         2         1     3         2         4
## 9         0           0       0   0   2         0         0     2         0         3
## 10        0           1       0   0   4         1         1     1         0         4
##      score
## 1         5
## 2         1
## 3         8
## 4         0
## 5         4
## 6         2
## 7         1
## 8         2
## 9         3
## 10        4
```

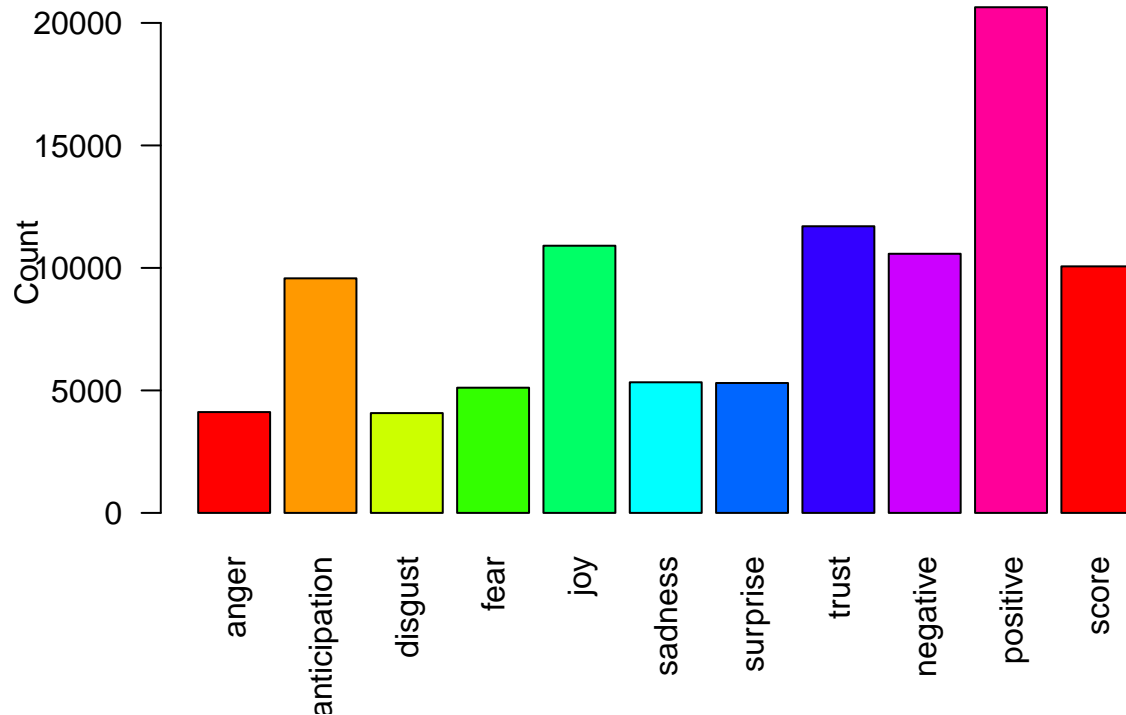
```
#check overall sentiment and represent it visually with a bar plot.
review_score <- colSums(s[,])
print(review_score)
```

```
##      anger anticipation      disgust      fear      joy      sadness
##      4115      9574      4074      5110      10905      5331
##      surprise      trust      negative      positive      score
##      5302      11700      10577      20640      10063
```

```
#Bar plot
barplot(colSums(s),
        las = 2,
        col = rainbow(10),
        ylab = 'Count',
        main = 'Sentiment')
```

productid	negative	positive	sentiment
141278509X	0	1	1
2734888454	2	4	2
2841233731	0	3	3
6641040	32	68	36
7310172001	144	143	-1
7310172101	144	143	-1

Sentiment



```
#sentiment score by productid
#Next, we can look at the sentiment score by product id and also see a visual
#representation of how top words contribute to "negative" and "positive"
#emotions.
```

```
sentiment_data <- food_reviews %>%
  inner_join(get_sentiments("bing"), "word") %>%
  count(productid, sentiment) %>%
  spread(sentiment, n, fill = 0) %>%
  mutate(sentiment = positive - negative)

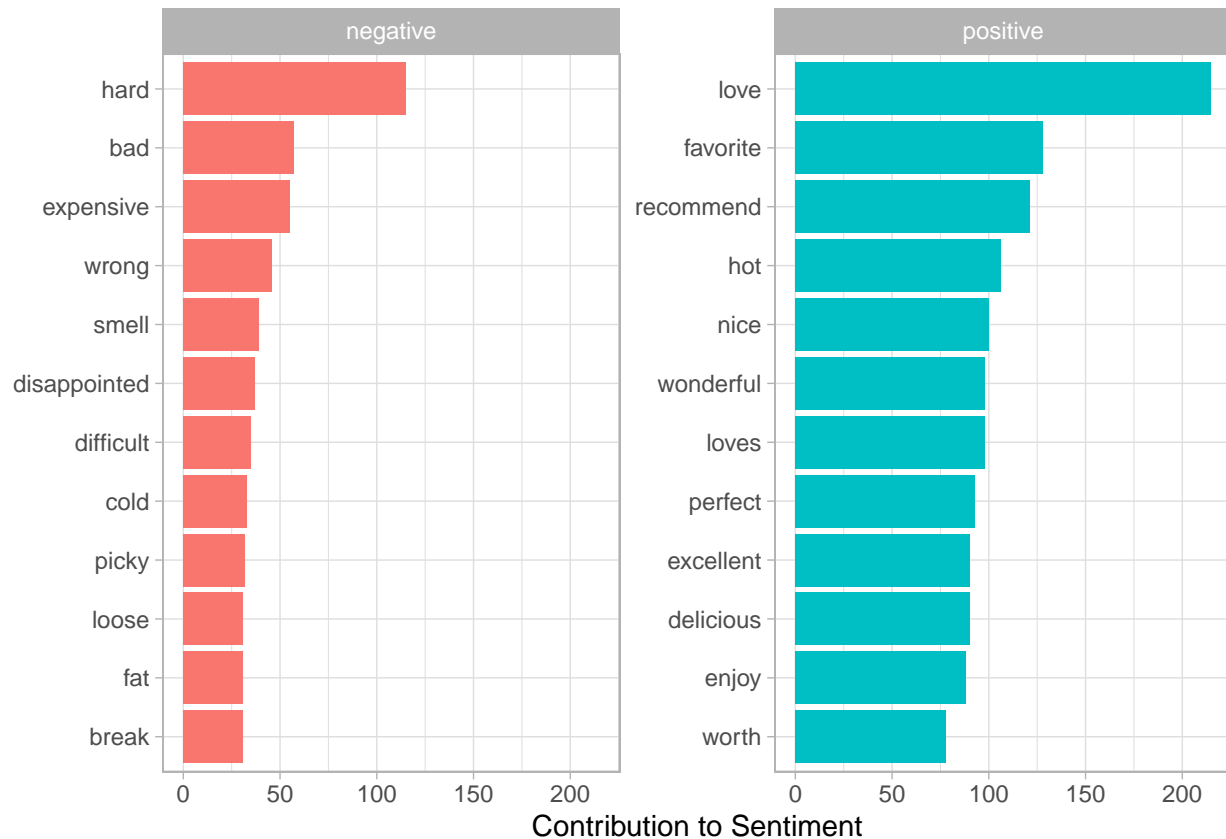
head(sentiment_data)%>%
  kable() %>%
  kable_styling(bootstrap_options = "basic", full_width = F)
```

```
sentiment_analysis_word_count <- food_reviews %>%
  inner_join(get_sentiments("bing"), "word") %>%
  count(word, sentiment, score, sort = TRUE) %>%
  ungroup()
```

```

sentiment_analysis_word_count %>%
  group_by(sentiment) %>%
  top_n(12, n) %>%
  ungroup() %>%
  mutate(word = reorder(word, n)) %>%
  ggplot(aes(word, n, fill = sentiment)) +
  geom_col(show.legend = FALSE) +
  facet_wrap(~sentiment, scales = "free_y") +
  labs(y = "Contribution to Sentiment", x = NULL) +
  coord_flip()

```



*#Finally, I looked at the avg. score (as provided by the initial review dataset)
 #to see whether there was a significant difference between
 #"positive" and "negative" sentiments. As expected, the "negative" category
 #was lower than the "positive" category, but by a small margin.*

```

#Avg. customer score by sentiment
sentiment_analysis_word_count %>%
  group_by(sentiment) %>%
  summarise_at(vars(score), list(name = mean))

```

```

## # A tibble: 2 x 2
##   sentiment name
##   <chr>      <dbl>
## 1 negative  3.42
## 2 positive  3.67

```